

1. Getting the data files in dataframe

```
1 import pandas as pd
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 train = pd.read_csv('https://raw.githubusercontent.com/kuldeep27396/AVH/main/train-file.csv')
5 train.head()
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product
0	NNVBBKZB	Female	73	RG268	Other	X3	43	No
1	IDD62UNG	Female	30	RG277	Salaried	X1	32	No
2	HD3DSEMC	Female	56	RG268	Self_Employed	X3	26	No
3	BF3NC7KV	Male	34	RG270	Salaried	X1	19	No
4	TEASRWXV	Female	30	RG282	Salaried	X1	33	No

```
1 test = pd.read_csv('https://raw.githubusercontent.com/kuldeep27396/AVH/main/test-file.csv')
2 test.head()
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Av
0	VBENBARO	Male	29	RG254	Other	X1	25	Yes	
1	CCMEWNKY	Male	43	RG268	Other	X2	49	NaN	
2	VK3KGA9M	Male	31	RG270	Salaried	X1	14	No	
3	TT8RPZVC	Male	29	RG272	Other	X1	33	No	
4	SHQZEY TZ	Female	29	RG270	Other	X1	19	No	

```
1 #copied train df to ff
2 ff= train.copy()
```

```
1 #copied test df to ff
2 gg = test.copy()
```

EDA


```
1 sns.countplot(train['Occupation']).tick_params(axis='x', rotation = 90)
2 plt.title('Occupation')
3 plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'Age'}. This will ensure compatibility in future versions of seaborn.



```
1 sns.set(rc={'figure.figsize':(12,5)})
2 sns.countplot(ff['Age']).tick_params(axis='x', rotation = 90)
3 plt.title('Age')
4 plt.show()
```


/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'Age'}. This will ensure compatibility in future versions of seaborn.



```
1 # To check the correlation among variables
2 plt.figure(figsize=(10,5))
3 sns.heatmap(ff.corr())
4 plt.show()
```



```
1 ax = sns.boxplot(data=ff, orient="h", palette="Set2")
```



2. Train Data Preprocessing

- Converted **Credit_Product**, **Gender**, **Is_Active** to binary
- Encoded **Occupation**, **Channel_code** to numeric
- **Dropped ID** from train as it has no corelation with our target

```
1 import numpy as np
2 ff.replace(to_replace = np.nan, value =5, inplace =True) #replaced all Nan values to 5(selected
3 ff['Credit_Product']=np.where(ff['Credit_Product']=='No',0,ff['Credit_Product'])
4 ff['Credit_Product']=np.where(ff['Credit_Product']=='Yes',1,ff['Credit_Product'])
5
6 ff['Is_Active'] = ff['Is_Active'].map({'Yes': 1, 'No': 0})
7 ff['Gender'] = ff['Gender'].map({'Male': 1, 'Female': 0})
8 ff['Occupation'] = ff['Occupation'].map({'Other': 0, 'Salaried': 1, 'Self_Employed': 2, 'Entrep
9 ff['Channel_Code'] = ff['Channel_Code'].map({'X1': 0, 'X2': 1, 'X3': 2, 'X4': 3})
10 ff.drop(['ID'],axis =1, inplace = True)
```

Encode labels in column 'Region_Code'

```
1 # Import label encoder
2 from sklearn import preprocessing
```

```

3
4 # label_encoder object knows how to understand word labels.
5 label_encoder = preprocessing.LabelEncoder()
6
7 # Encode labels in column 'Region_Code'.
8 ff['Region_Code']= label_encoder.fit_transform(ff['Region_Code'])
9
10 ff['Region_Code'].unique()
11
array([18, 27, 20, 32, 11, 15, 33, 4, 19, 7, 29, 30, 2, 34, 9, 31, 8,
       16, 10, 24, 6, 25, 23, 17, 22, 1, 12, 14, 28, 26, 13, 0, 5, 3,
       21])

```

1

Avg_Account_Balance have outliers as shown in the image, I have done label encoding and also removed outliers as shown below.

Encode labels in column 'Avg_Account_Balance'

```

1 # Import label encoder
2 from sklearn import preprocessing
3
4 # label_encoder object knows how to understand word labels.
5 label_encoder = preprocessing.LabelEncoder()
6
7 # Encode labels in column 'Avg_Account_Balance'.
8 ff['Avg_Account_Balance']= label_encoder.fit_transform(ff['Avg_Account_Balance'])
9
10 ff['Avg_Account_Balance'].unique()

array([ 69898, 28951, 96873, ..., 9819, 112817, 76068])

```

```

1 #import re
2 #replace = re.compile("([a-zA-Z]+)")
3 #ff['Region_Code'] = ff['Region_Code'].astype(str).replace(replace, "")

```

```

1 #get the dataframe after transformation
2 ff.head()

```

	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Avg_Account_Ba
0	0	73	18	0	2	43	0	
1	0	30	27	1	0	32	0	
2	0	56	18	2	2	26	0	
3	1	34	20	1	0	19	0	
4	0	30	32	1	0	33	0	

Removing outliers from Avg_Account_Balance

```

1  # Importing
2  import sklearn
3  import pandas as pd
4
5  # IQR
6  Q1 = np.percentile(ff['Avg_Account_Balance'], 25,
7                      interpolation = 'midpoint')
8
9  Q3 = np.percentile(ff['Avg_Account_Balance'], 75,
10                     interpolation = 'midpoint')
11  IQR = Q3 - Q1
12
13  print("Old Shape: ", ff.shape)
14
15  # Upper bound
16  upper = np.where(ff['Avg_Account_Balance'] >= (Q3+1.5*IQR))
17  # Lower bound
18  lower = np.where(ff['Avg_Account_Balance'] <= (Q1-1.5*IQR))
19
20  ''' Removing the Outliers '''
21  ff.drop(upper[0], inplace = True)
22  ff.drop(lower[0], inplace = True)
23
24  print("New Shape: ", ff.shape)
25

```


Old Shape: (245725, 10)

New Shape: (245725, 10)

```

1  ax = sns.boxplot(data=ff, orient="h", palette="Set2")

```



```

1  ff['Credit_Product'].value_counts()

```

0 144357

1 72043

5 29325

Name: Credit_Product, dtype: int64

3. Test data processing

Everything will be same as train preprocessing

```
1 gg.head()
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Av
0	VBENBARO	Male	29	RG254	Other	X1	25	Yes	
1	CCMEWNKY	Male	43	RG268	Other	X2	49	NaN	
2	VK3KGA9M	Male	31	RG270	Salaried	X1	14	No	
3	TT8RPZVC	Male	29	RG272	Other	X1	33	No	
4	SHQZEYtz	Female	29	RG270	Other	X1	19	No	

```
1 gg.replace(to_replace = np.nan, value =5, inplace =True)
2 gg['Credit_Product']=np.where(gg['Credit_Product']=='No',0,gg['Credit_Product'])
3 gg['Credit_Product']=np.where(gg['Credit_Product']=='Yes',1,gg['Credit_Product'])
4
5 gg['Channel_Code'] = gg['Channel_Code'].map({'X1': 0, 'X2': 1, 'X3': 2, 'X4': 3})
6
7 gg['Gender'] = gg['Gender'].map({'Male': 1, 'Female': 0})
8
9 gg['Occupation'] = gg['Occupation'].map({'Other': 0, 'Salaried': 1, 'Self_Employed': 2, 'Entrepreneur': 3})
10 gg['Is_Active'] = gg['Is_Active'].map({'Yes': 1, 'No': 0})
```

```
1 # Import label encoder
2 from sklearn import preprocessing
3
4 # label_encoder object knows how to understand word labels.
5 label_encoder = preprocessing.LabelEncoder()
6
7 # Encode labels in column 'Avg_Account_Balance'.
8 gg['Avg_Account_Balance']= label_encoder.fit_transform(gg['Avg_Account_Balance'])
9
10 gg['Avg_Account_Balance'].unique()
```

```
array([27635, 37999, 732, ..., 48831, 63962, 18118])
```

```
1 # Import label encoder
2 from sklearn import preprocessing
3
4 # label_encoder object knows how to understand word labels.
5 label_encoder = preprocessing.LabelEncoder()
6
7 # Encode labels in column 'Region_Code'.
8 gg['Region_Code']= label_encoder.fit_transform(gg['Region_Code'])
9
10 gg['Region_Code'].unique()
```

```
array([ 4, 18, 20, 22,  3,  7, 34, 33, 27, 23, 11, 30, 24, 28, 19, 31,  1,
        2, 29, 32, 12, 13, 14,  6, 15, 10,  8, 21,  0, 25, 26, 17,  5,  9,
        16])
```

```
1 gg['Is_Active'].value_counts()
```

```
0    63797
1    41515
Name: Is_Active, dtype: int64
```

4. Test-Train Split

```
1 # Import the required library
2 from sklearn.model_selection import train_test_split
```

```
1 X = ff.drop(['Is_Lead'], 1)
2 X.head()
```

	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Avg_Account_Bal
0	0	73	18	0	2	43	0	
1	0	30	27	1	0	32	0	
2	0	56	18	2	2	26	0	
3	1	34	20	1	0	19	0	
4	0	30	32	1	0	33	0	

```
1 # Putting the target variable in y
2 y = ff['Is_Lead']
3 y.head()
```

```
0    0
1    0
2    0
3    0
4    0
Name: Is_Lead, dtype: int64
```

```
1 # Split the dataset into 70% and 30% for train and test respectively(taken with imp features)
2 X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size=0.3, random=
```

```
1 X_train['Is_Active'].value_counts()
```

```
0    105293
1     66714
Name: Is_Active, dtype: int64
```

5. Model Building : Logistic Regression

```
1 #Let's start with importing necessary libraries
2 import pandas as pd
3 import numpy as np
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.linear_model import Ridge,Lasso,RidgeCV, LassoCV, ElasticNet, ElasticNetCV, LogisticRegression
6 from sklearn.model_selection import train_test_split
```

```
7 from statsmodels.stats.outliers_influence import variance_inflation_factor
8 from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc_score
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 sns.set()
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas
import pandas.util.testing as tm
```

```
1 log_reg = LogisticRegression()
2 log_reg.fit(X_train,y_train)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

Feature selection with RFE

```
1 # Running RFE with 6 variables as output
2 from sklearn.feature_selection import RFE
3 rfe = RFE(log_reg,6)
4 rfe = rfe.fit(X_train, y_train)
```

```
1 rfe.support_
```

```
array([ True, False, False,  True,  True,  True,  True, False,  True])
```

```
1 # Features that have been selected by RFE
2 list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[('Gender', True, 1),
 ('Age', False, 2),
 ('Region_Code', False, 3),
 ('Occupation', True, 1),
 ('Channel_Code', True, 1),
 ('Vintage', True, 1),
 ('Credit_Product', True, 1),
 ('Avg_Account_Balance', False, 4),
 ('Is_Active', True, 1)]
```

```
1 # Put all the columns selected by RFE in the variable 'col'
2 col = X_train.columns[rfe.support_]
```

```
1 col
```

```
Index(['Gender', 'Occupation', 'Channel_Code', 'Vintage', 'Credit_Product',
      'Is_Active'],
      dtype='object')
```

```
1 # Selecting columns selected by RFE
2 X_train1 = X_train[col]
```


Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: plotly in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.7/dist-packages (from CatBoost)

Installing collected packages: CatBoost

Successfully installed CatBoost-0.25.1

```
1 #with all features
2 model2 = CatBoostClassifier(iterations=100)
3
4 # fit the model with the training data
5 model2.fit(X_train,y_train,plot=False)
6 print('\n Model Trained')
7
8 # predict the target on the train dataset
9 predict_train = model2.predict(X_train)
10
```

42:	learn: 0.3390712	total: 1.47s	remaining: 1.94s
43:	learn: 0.3389433	total: 1.5s	remaining: 1.91s
44:	learn: 0.3388630	total: 1.53s	remaining: 1.87s
45:	learn: 0.3388316	total: 1.56s	remaining: 1.83s
46:	learn: 0.3387661	total: 1.59s	remaining: 1.79s
47:	learn: 0.3386568	total: 1.62s	remaining: 1.76s
48:	learn: 0.3385286	total: 1.65s	remaining: 1.72s
49:	learn: 0.3384583	total: 1.7s	remaining: 1.7s
50:	learn: 0.3383273	total: 1.73s	remaining: 1.66s
51:	learn: 0.3382499	total: 1.77s	remaining: 1.63s
52:	learn: 0.3381685	total: 1.8s	remaining: 1.59s
53:	learn: 0.3381484	total: 1.83s	remaining: 1.56s
54:	learn: 0.3380283	total: 1.86s	remaining: 1.52s
55:	learn: 0.3379584	total: 1.89s	remaining: 1.49s
56:	learn: 0.3378348	total: 1.93s	remaining: 1.45s
57:	learn: 0.3377565	total: 1.96s	remaining: 1.42s
58:	learn: 0.3376638	total: 1.99s	remaining: 1.39s
59:	learn: 0.3375980	total: 2.03s	remaining: 1.35s
60:	learn: 0.3375130	total: 2.06s	remaining: 1.31s
61:	learn: 0.3374044	total: 2.09s	remaining: 1.28s
62:	learn: 0.3373924	total: 2.12s	remaining: 1.24s
63:	learn: 0.3373058	total: 2.15s	remaining: 1.21s
64:	learn: 0.3372280	total: 2.18s	remaining: 1.17s
65:	learn: 0.3371683	total: 2.21s	remaining: 1.14s
66:	learn: 0.3370851	total: 2.24s	remaining: 1.1s
67:	learn: 0.3369372	total: 2.28s	remaining: 1.07s
68:	learn: 0.3368808	total: 2.32s	remaining: 1.04s
69:	learn: 0.3367826	total: 2.35s	remaining: 1.01s
70:	learn: 0.3366921	total: 2.38s	remaining: 974ms
71:	learn: 0.3365919	total: 2.43s	remaining: 944ms
72:	learn: 0.3364775	total: 2.47s	remaining: 912ms
73:	learn: 0.3363537	total: 2.5s	remaining: 878ms
74:	learn: 0.3362945	total: 2.53s	remaining: 843ms
75:	learn: 0.3361972	total: 2.56s	remaining: 809ms
76:	learn: 0.3361174	total: 2.59s	remaining: 774ms
77:	learn: 0.3360278	total: 2.62s	remaining: 740ms
78:	learn: 0.3359646	total: 2.66s	remaining: 706ms
79:	learn: 0.3359204	total: 2.69s	remaining: 672ms

80:	learn: 0.3358490	total: 2.72s	remaining: 638ms
81:	learn: 0.3357755	total: 2.75s	remaining: 604ms
82:	learn: 0.3357096	total: 2.78s	remaining: 569ms
83:	learn: 0.3356039	total: 2.81s	remaining: 535ms
84:	learn: 0.3355404	total: 2.84s	remaining: 501ms
85:	learn: 0.3354675	total: 2.88s	remaining: 469ms
86:	learn: 0.3354332	total: 2.91s	remaining: 435ms
87:	learn: 0.3353857	total: 2.94s	remaining: 401ms
88:	learn: 0.3353028	total: 2.97s	remaining: 367ms
89:	learn: 0.3352315	total: 3s	remaining: 334ms
90:	learn: 0.3351515	total: 3.03s	remaining: 300ms
91:	learn: 0.3350505	total: 3.06s	remaining: 267ms
92:	learn: 0.3350037	total: 3.1s	remaining: 233ms
93:	learn: 0.3349048	total: 3.13s	remaining: 200ms
94:	learn: 0.3348624	total: 3.16s	remaining: 166ms
95:	learn: 0.3347966	total: 3.19s	remaining: 133ms
96:	learn: 0.3347316	total: 3.22s	remaining: 99.7ms
97:	learn: 0.3346657	total: 3.25s	remaining: 66.4ms
98:	learn: 0.3346313	total: 3.28s	remaining: 33.2ms
99:	learn: 0.3345538	total: 3.32s	remaining: 0us

Model Trained

```

1 # Running RFE with 4 variables as output
2 from sklearn.feature_selection import RFE
3 rfe = RFE(model2,7)
4 rfe = rfe.fit(X_train, y_train)

```

40:	learn: 0.3402027	total: 1.29s	remaining: 1.86s
41:	learn: 0.3401493	total: 1.32s	remaining: 1.82s
42:	learn: 0.3399346	total: 1.35s	remaining: 1.79s
43:	learn: 0.3398523	total: 1.38s	remaining: 1.76s
44:	learn: 0.3397789	total: 1.41s	remaining: 1.73s
45:	learn: 0.3397065	total: 1.44s	remaining: 1.69s
46:	learn: 0.3395877	total: 1.48s	remaining: 1.66s
47:	learn: 0.3395205	total: 1.51s	remaining: 1.64s
48:	learn: 0.3394008	total: 1.54s	remaining: 1.6s
49:	learn: 0.3393227	total: 1.57s	remaining: 1.57s
50:	learn: 0.3392520	total: 1.6s	remaining: 1.54s
51:	learn: 0.3392015	total: 1.63s	remaining: 1.5s
52:	learn: 0.3391566	total: 1.66s	remaining: 1.47s
53:	learn: 0.3391031	total: 1.69s	remaining: 1.44s
54:	learn: 0.3389804	total: 1.72s	remaining: 1.41s
55:	learn: 0.3388843	total: 1.75s	remaining: 1.38s
56:	learn: 0.3386989	total: 1.78s	remaining: 1.35s
57:	learn: 0.3385967	total: 1.82s	remaining: 1.31s
58:	learn: 0.3385083	total: 1.85s	remaining: 1.28s
59:	learn: 0.3384307	total: 1.89s	remaining: 1.26s
60:	learn: 0.3383233	total: 1.92s	remaining: 1.23s
61:	learn: 0.3382443	total: 1.95s	remaining: 1.2s
62:	learn: 0.3381674	total: 1.98s	remaining: 1.16s
63:	learn: 0.3381157	total: 2.01s	remaining: 1.13s
64:	learn: 0.3380388	total: 2.04s	remaining: 1.1s
65:	learn: 0.3379332	total: 2.07s	remaining: 1.07s
66:	learn: 0.3378293	total: 2.1s	remaining: 1.04s
67:	learn: 0.3377673	total: 2.14s	remaining: 1.01s
68:	learn: 0.3377045	total: 2.18s	remaining: 978ms
69:	learn: 0.3376574	total: 2.22s	remaining: 953ms
70:	learn: 0.3375585	total: 2.25s	remaining: 921ms
71:	learn: 0.3374805	total: 2.28s	remaining: 888ms
72:	learn: 0.3374344	total: 2.31s	remaining: 856ms
73:	learn: 0.3374093	total: 2.34s	remaining: 823ms
74:	learn: 0.3373401	total: 2.37s	remaining: 791ms
75:	learn: 0.3372891	total: 2.41s	remaining: 760ms
76:	learn: 0.3372191	total: 2.44s	remaining: 728ms
77:	learn: 0.3371459	total: 2.47s	remaining: 697ms

78:	learn: 0.3371133	total: 2.5s	remaining: 665ms
79:	learn: 0.3370630	total: 2.53s	remaining: 633ms
80:	learn: 0.3370060	total: 2.56s	remaining: 600ms
81:	learn: 0.3369070	total: 2.59s	remaining: 569ms
82:	learn: 0.3368732	total: 2.62s	remaining: 537ms
83:	learn: 0.3368452	total: 2.65s	remaining: 505ms
84:	learn: 0.3367455	total: 2.69s	remaining: 474ms
85:	learn: 0.3367153	total: 2.72s	remaining: 442ms
86:	learn: 0.3366887	total: 2.75s	remaining: 410ms
87:	learn: 0.3366336	total: 2.77s	remaining: 378ms
88:	learn: 0.3365712	total: 2.8s	remaining: 347ms
89:	learn: 0.3365343	total: 2.84s	remaining: 316ms
90:	learn: 0.3364687	total: 2.87s	remaining: 284ms
91:	learn: 0.3364084	total: 2.9s	remaining: 252ms
92:	learn: 0.3363835	total: 2.93s	remaining: 220ms
93:	learn: 0.3363709	total: 2.96s	remaining: 189ms
94:	learn: 0.3363109	total: 2.99s	remaining: 157ms
95:	learn: 0.3362272	total: 3.02s	remaining: 126ms
96:	learn: 0.3361722	total: 3.05s	remaining: 94.4ms
97:	learn: 0.3361373	total: 3.08s	remaining: 62.9ms
98:	learn: 0.3360994	total: 3.11s	remaining: 31.4ms
99:	learn: 0.3360231	total: 3.15s	remaining: 0us

```
1 rfe.support_
```

```
array([False,  True, False,  True,  True,  True,  True,  True,  True])
```

```
1 # Features that have been selected by RFE
```

```
2 list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[('Gender', False, 3),
 ('Age', True, 1),
 ('Region_Code', False, 2),
 ('Occupation', True, 1),
 ('Channel_Code', True, 1),
 ('Vintage', True, 1),
 ('Credit_Product', True, 1),
 ('Avg_Account_Balance', True, 1),
 ('Is_Active', True, 1)]
```

```
1 # Put all the columns selected by RFE in the variable 'col'
```

```
2 col = X_train.columns[rfe.support_]
```

```
1 X_train11 = X_train[col]
```

```
1 #with all features
```

```
2 model2 = CatBoostClassifier(iterations=100)
```

```
3
```

```
4 # fit the model with the training data
```

```
5 model2.fit(X_train11,y_train,plot=False)
```

```
6 print('\n Model Trained')
```

```
7
```

```
8 # predict the target on the train dataset
```

```
9 predict_train = model2.predict(X_train11)
```

42:	learn: 0.3399346	total: 1.36s	remaining: 1.8s
43:	learn: 0.3398523	total: 1.39s	remaining: 1.77s
44:	learn: 0.3397789	total: 1.42s	remaining: 1.74s
45:	learn: 0.3397065	total: 1.46s	remaining: 1.71s
46:	learn: 0.3395877	total: 1.49s	remaining: 1.68s
47:	learn: 0.3395205	total: 1.52s	remaining: 1.65s

48:	learn: 0.3394008	total: 1.55s	remaining: 1.62s
49:	learn: 0.3393227	total: 1.58s	remaining: 1.58s
50:	learn: 0.3392520	total: 1.61s	remaining: 1.55s
51:	learn: 0.3392015	total: 1.64s	remaining: 1.51s
52:	learn: 0.3391566	total: 1.67s	remaining: 1.48s
53:	learn: 0.3391031	total: 1.7s	remaining: 1.45s
54:	learn: 0.3389804	total: 1.75s	remaining: 1.43s
55:	learn: 0.3388843	total: 1.78s	remaining: 1.4s
56:	learn: 0.3386989	total: 1.82s	remaining: 1.37s
57:	learn: 0.3385967	total: 1.85s	remaining: 1.34s
58:	learn: 0.3385083	total: 1.88s	remaining: 1.31s
59:	learn: 0.3384307	total: 1.91s	remaining: 1.27s
60:	learn: 0.3383233	total: 1.94s	remaining: 1.24s
61:	learn: 0.3382443	total: 1.98s	remaining: 1.21s
62:	learn: 0.3381674	total: 2.01s	remaining: 1.18s
63:	learn: 0.3381157	total: 2.04s	remaining: 1.15s
64:	learn: 0.3380388	total: 2.08s	remaining: 1.12s
65:	learn: 0.3379332	total: 2.1s	remaining: 1.08s
66:	learn: 0.3378293	total: 2.14s	remaining: 1.05s
67:	learn: 0.3377673	total: 2.17s	remaining: 1.02s
68:	learn: 0.3377045	total: 2.2s	remaining: 990ms
69:	learn: 0.3376574	total: 2.23s	remaining: 957ms
70:	learn: 0.3375585	total: 2.26s	remaining: 924ms
71:	learn: 0.3374805	total: 2.29s	remaining: 891ms
72:	learn: 0.3374344	total: 2.32s	remaining: 859ms
73:	learn: 0.3374093	total: 2.35s	remaining: 826ms
74:	learn: 0.3373401	total: 2.38s	remaining: 794ms
75:	learn: 0.3372891	total: 2.42s	remaining: 763ms
76:	learn: 0.3372191	total: 2.45s	remaining: 731ms
77:	learn: 0.3371459	total: 2.48s	remaining: 699ms
78:	learn: 0.3371133	total: 2.5s	remaining: 666ms
79:	learn: 0.3370630	total: 2.54s	remaining: 634ms
80:	learn: 0.3370060	total: 2.56s	remaining: 602ms
81:	learn: 0.3369070	total: 2.6s	remaining: 570ms
82:	learn: 0.3368732	total: 2.63s	remaining: 539ms
83:	learn: 0.3368452	total: 2.66s	remaining: 507ms
84:	learn: 0.3367455	total: 2.7s	remaining: 476ms
85:	learn: 0.3367153	total: 2.73s	remaining: 444ms
86:	learn: 0.3366887	total: 2.76s	remaining: 413ms
87:	learn: 0.3366336	total: 2.79s	remaining: 381ms
88:	learn: 0.3365712	total: 2.82s	remaining: 349ms
89:	learn: 0.3365343	total: 2.86s	remaining: 318ms
90:	learn: 0.3364687	total: 2.89s	remaining: 286ms
91:	learn: 0.3364084	total: 2.92s	remaining: 254ms
92:	learn: 0.3363835	total: 2.96s	remaining: 223ms
93:	learn: 0.3363709	total: 2.99s	remaining: 191ms
94:	learn: 0.3363109	total: 3.02s	remaining: 159ms
95:	learn: 0.3362272	total: 3.05s	remaining: 127ms
96:	learn: 0.3361722	total: 3.09s	remaining: 95.5ms
97:	learn: 0.3361373	total: 3.12s	remaining: 63.6ms
98:	learn: 0.3360994	total: 3.14s	remaining: 31.8ms
99:	learn: 0.3360231	total: 3.17s	remaining: 0us

Model Trained

```
1 X_test11 = X_test[col]
```

```
1 y_pred11 = model2.predict(X_test11)
2 auc = roc_auc_score(y_test, y_pred11)
3 auc
```

0.7512887445265599

7. Model Building: ANN Classification

```

1  #converting to scaler
2  from sklearn.preprocessing import StandardScaler
3  sc = StandardScaler()
4  X_train = sc.fit_transform(X_train)
5  X_test = sc.transform(X_test)

```

```

1  from keras.models import Sequential
2  classifier = Sequential()

```

```

1  from keras.models import Sequential
2  from keras.layers import Dense
3  from keras.wrappers.scikit_learn import KerasClassifier
4  from sklearn.model_selection import cross_val_score
5  from sklearn.preprocessing import LabelEncoder
6  from sklearn.model_selection import StratifiedKFold
7  from sklearn.preprocessing import StandardScaler
8  from sklearn.pipeline import Pipeline

```

```

1  #Initializing Neural Network
2  from keras.models import Sequential
3  from keras.layers import Activation, Dense, Dropout
4  from keras import optimizers
5  Model = Sequential()
6  Model.add(Dense(64, kernel_initializer = 'uniform', activation = 'relu', input_dim = 9))
7  # Adding the second hidden layer
8  Model.add(Dense(48, kernel_initializer = 'uniform', activation = 'relu'))
9  Model.add(Dense(48, kernel_initializer = 'uniform', activation = 'relu'))
10 Model.add(Dense(48, kernel_initializer = 'uniform', activation = 'relu'))
11 Model.add(Dense(32, kernel_initializer = 'uniform', activation = 'tanh'))
12 Model.add(Dense(32, kernel_initializer = 'uniform', activation = 'relu'))
13 Model.add(Dropout(0.2))
14 # Adding the output layer
15 Model.add(Dense(1, kernel_initializer = 'uniform', activation = 'sigmoid'))
16 # Compiling Neural Network
17 Model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['AUC'],)
18 #fitting the neural Network
19 Model.fit(X_train, y_train, batch_size = 48, epochs = 30)

```

```

Epoch 2/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3631 - auc: 0.8579
Epoch 3/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3553 - auc: 0.8625
Epoch 4/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3498 - auc: 0.8676
Epoch 5/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3471 - auc: 0.8681
Epoch 6/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3500 - auc: 0.8654
Epoch 7/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3508 - auc: 0.8671
Epoch 8/30

```

```

3584/3584 [=====] - 7s 2ms/step - loss: 0.3476 - auc: 0.8693
Epoch 9/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3484 - auc: 0.8686
Epoch 10/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3467 - auc: 0.8706
Epoch 11/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3441 - auc: 0.8713
Epoch 12/30
3584/3584 [=====] - 8s 2ms/step - loss: 0.3445 - auc: 0.8727
Epoch 13/30
3584/3584 [=====] - 8s 2ms/step - loss: 0.3456 - auc: 0.8707
Epoch 14/30
3584/3584 [=====] - 8s 2ms/step - loss: 0.3450 - auc: 0.8719
Epoch 15/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3436 - auc: 0.8715
Epoch 16/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3460 - auc: 0.8720
Epoch 17/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3468 - auc: 0.8706
Epoch 18/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3437 - auc: 0.8728
Epoch 19/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3466 - auc: 0.8717
Epoch 20/30
3584/3584 [=====] - 8s 2ms/step - loss: 0.3434 - auc: 0.8730
Epoch 21/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3465 - auc: 0.8692
Epoch 22/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3445 - auc: 0.8726
Epoch 23/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3439 - auc: 0.8728
Epoch 24/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3455 - auc: 0.8717
Epoch 25/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3459 - auc: 0.8713
Epoch 26/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3402 - auc: 0.8763
Epoch 27/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3461 - auc: 0.8711
Epoch 28/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3439 - auc: 0.8729
Epoch 29/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3432 - auc: 0.8737
Epoch 30/30
3584/3584 [=====] - 7s 2ms/step - loss: 0.3442 - auc: 0.8729
<keras.callbacks.History at 0x7fc24d8f0ad0>

```

Getting AUC

```

1 Y_prediction=Model.predict(X_test)
2 # Area Under Curve
3 auc = roc_auc_score(y_test, Y_prediction)
4 auc

```

```
0.8713652724144249
```

Got final AUC:0.8721

1. Got the testing data ID in ID column

2. Dropped ID from testing data to do after preprocessing i.e. converting to Scaler

```
1 ID = gg['ID']
2 gg.drop(['ID'], axis =1, inplace =True)
```

Converting to Scaler the testing data

```
1 from sklearn.preprocessing import StandardScaler
2 sc = StandardScaler()
3 X_train55 = sc.fit_transform(gg)
```

```
1 X_train55
```

```
array([[ 0.90829868, -0.99987782, -1.45176439, ...,  0.06952535,
        -0.45614794, -0.80668205],
       [ 0.90829868, -0.05856336, -0.16779796, ...,  2.61183556,
         0.0167183 , -0.80668205],
       [ 0.90829868, -0.86540432,  0.01562582, ..., -0.56605221,
        -1.68362001, -0.80668205],
       ...,
       [ 0.90829868, -0.59645734, -1.45176439, ..., -0.56605221,
         1.20130207, -0.80668205],
       [ 0.90829868,  0.6138041 , -1.45176439, ..., -0.56605221,
        -0.47010945,  1.23964578],
       [ 0.90829868, -1.13435131, -1.26834061, ..., -0.56605221,
        -0.8903691 , -0.80668205]])
```

Generating the submission file

```
1 y_pred = Model.predict(X_train55)
2 y_pred = (y_pred > 0.25)
3 submission_df = pd.DataFrame({'ID': ID, 'Is_Lead': y_pred.flatten().astype(int),
4 })
5 submission_df.set_index('ID', inplace=True)
6 submission_df.to_csv('final_submission.csv')
```